

Associative Arrays

The Map ADT

Introduction

- In Java, there are three basic collection interfaces:
 1. The Set ADT: Stores an *unordered* collection of elements.
 2. The List ADT: Stores a position *ordered* collection of elements.
 3. The Map ADT: Stores a collection of *key-value* pair entries.
- In an ordinary array, a value stored in the array is directly accessed by an *integer index* corresponding to the *location* where that value is stored.
- A **Map** (some times called *associative array*) is an abstract generalization of the ordinary array concept so that a value stored in the array is directly accessed by its *associated key*.

The Map ADT Definition

- A **map** models a searchable collection of **key-value entries**, where each **key**, **k**, is **associated** with a corresponding **value**, **v**.
- The main **operations** of a map are for **searching**, **inserting**, and **deleting entries**.
- Multiple entries with the same **key** are **not** allowed.
- There are two types of maps: **Sorted** or **Unsorted**.
- The **keys** and **values** can be of **any** object type.
- **Example Applications**:
 - Address book.
 - Student record database.
 - Compiler's symbol table.

The Unsorted Map ADT Operations

- **Unsorted Map ADT Operations**:
 - **get(k)**: if the map **M** has an **entry** with key **k**, return its associated value, **v**; else, return **null**.
 - **put(k, v)**: insert **entry** (**k**, **v**) into the map **M**, keeping **k** unique; if key **k** is not already in **M**, then return **null**; else, replace and return the old value, **v** associated with **k**.
 - **remove(k)**: if the map **M** has an **entry** with key **k**, remove it from **M** and return its associated value, **v**; else, return **null**.
 - **size()**, **isEmpty()**: As before.
 - **keys()**: return an **iterable** collection of all the keys in **M**.
 - **values()**: return an **iterable** collection of the values of all **entries** in **M**.
 - **entries()**: return an **iterable** collection of all the key-value **entries** in **M**.

Unsorted Map Implementations

- An **Unsorted Map** can be implemented using four main data structures:
 - An **unordered list**, (e.g. **ArrayList** or **LinkedList**).
 - A **search table**, (e.g. Sorted **ArrayList**).
 - A **hash table**, using bucket array & external chaining.
 - A **skip list**.

Comparing Unsorted Map Implementations

- All implementations require $O(n)$ space. Assume the following:
 n : The number of entries in the Map,
- The time requirements of all operations are shown in the following table:

Operation Using	get	put	remove	entries	size	isEmpty	clear
Unordered list	$O(n)$	$O(n)$	$O(n)$	$O(n)$	$O(1)$	$O(1)$	$O(1)$
Search table	$O(\log n)$	$O(n)$	$O(n)$	$O(n)$	$O(1)$	$O(1)$	$O(1)$
Hash table	$O(1)$ exp. $O(n)$ worst	$O(1)$ exp. $O(n)$ worst	$O(1)$ exp. $O(n)$ worst	$O(n)$	$O(1)$	$O(1)$	$O(1)$
Skip List	$O(\log n)$ exp.	$O(\log n)$ exp.	$O(\log n)$ exp.	$O(n)$	$O(1)$	$O(1)$	$O(1)$

The Sorted Map ADT Operations

- Sorted Map ADT Operations:

- All the unsorted Map operations as before, and the following operations:
- first()**: Returns the **entry** with smallest key, or **null**, if the map is empty.
- last()**: Returns the **entry** with largest key, or **null**, if the map is empty.
- ceiling(k)**: Returns the **entry** with the least key greater than or equal to **k**, or **null**, if no such entry exists.
- floor(k)**: Returns the **entry** with the greatest key less than or equal to **k**, or **null**, if no such entry exists.
- lower(k)**: Returns the **entry** with the greatest key strictly less than **k**, or **null**, if no such entry exists.
- higher(k)**: Returns the **entry** with the least key strictly greater than **k**, or **null**, if no such entry exists.
- subMap(k₁, k₂)**: Returns an **iterable** collection of all the **entries** in **M** with keys greater than or equal to **k₁** and strictly less than **k₂**.

Sorted Map Implementations

- A **Sorted Map** can be implemented efficiently using three main data structures:
 - A **balanced search tree**, (e.g. **red-black** or **AVL binary trees**); and a **B-tree** (or its variants) can also be used when the map is too large to reside entirely in main memory.
 - A **search table**, (e.g. **Sorted ArrayList**).
 - A **skip list**.

Comparing Sorted Map Implementations

- All implementations require $O(n)$ space. Assume the following:
 - n : The number of entries in the sorted Map,
 - s : The size of the collection returned by the `subMap()` operation.
- The time requirements of all operations are shown in the following table:

Operation Using	subMap	get	put	remove	first	last	ceiling, higher	floor, lower	keys, values, entries
Balanced Search Tree	$O(\log n + s)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$		$O(\log n)$		$O(n)$
Search table	$O(\log n + s)$	$O(\log n)$	$O(n)$	$O(n)$	$O(1)$		$O(\log n)$		$O(n)$
Skip List	$O(\log n + s)$ exp.	$O(\log n)$ exp.	$O(\log n)$ exp.	$O(\log n)$ exp.	$O(1)$		$O(\log n)$ exp.		$O(n)$

Empty Slide

Dictionaries

The Dictionary ADT

The Dictionary ADT Definition

- A **dictionary**, like a **map**, models a searchable collection of **key-value entries**, where each **key**, **k**, is **associated** with a corresponding **value**, **v**.
 - Unlike a **map**, **Multiple entries with the same key are allowed**.
 - The **keys** and **values** can be of **any object** type.
- There are **two** dictionary types:
 - **Ordered**,
 - **Unordered**.
- The main **operations** of a dictionary are for **searching**, **inserting**, and **deleting entries**.
- Example Applications:
 - Log files or audit trails.
 - Language dictionaries.
 - Sorted sets.

Unordered Dictionary Operations

- Operations on an Unordered Dictionary D are:
 - **find**(k): if dictionary D has any entry with key k , it returns that entry; else, it returns **null**.
 - **insert**(k, v): inserts entry (k, v) into dictionary D , and returns the entry created.
 - **remove**(e): if dictionary D has entry e , it removes that entry from D and returns it; else, it returns **null**.
 - **size()**, **isEmpty()**, **clear()**: As defined before.
 - **findAll**(k): returns an **iterable** collection of all entries in D with keys = k .
 - **entries**(): returns an **iterable** collection of all the key-value entries in D .

Unordered Dictionary Implementations

- An **Unordered Dictionary** can be implemented efficiently using four main data structures:
 - An **unordered list**, (e.g. **ArrayList** or **LinkedList**).
 - A **search table**, (e.g. Sorted **ArrayList**).
 - A **hash table**, using bucket array & external chaining.
 - A **skip list**.

Comparing Unordered Dictionary Implementations

- All implementations require $O(n)$ space. Assume the following:
 - n : The number of entries in the dictionary,
 - s : The size of the collection returned by operation `findAll(k)`.
- The time requirements of all operations are shown in the following table:

Operation Using	find	findAll	remove	insert	entries	size	isEmpty	clear
Unordered list	$O(n)$	$O(n)$	$O(n)$	$O(1)$	$O(n)$	$O(1)$	$O(1)$	$O(1)$
Search table	$O(\log n)$	$O(\log n + s)$	$O(n)$	$O(n)$	$O(n)$	$O(1)$	$O(1)$	$O(1)$
Hash table	$O(1)$ exp. $O(n)$ worst	$O(1+s)$ exp. $O(n)$ worst	$O(1)$ exp. $O(n)$ worst	$O(1)$	$O(n)$	$O(1)$	$O(1)$	$O(1)$
Skip List	$O(\log n)$ exp.	$O(\log n + s)$ exp.	$O(\log n)$ exp.		$O(n)$	$O(1)$	$O(1)$	$O(1)$

Ordered Dictionary Operations

- Operations on an Ordered Dictionary D are:
 - All the operations of the unordered dictionary are the same for the ordered dictionary:
 - `find(k)`, `findAll(k)`, `entries()`: As defined before.
 - `insert(k, v)`, `remove(e)`: As defined before.
 - `size()`, `isEmpty()`, `clear()`: As defined before.
 - An ordered dictionary can efficiently provide the additional operations:
 - `first()`: returns an entry in D with the smallest key.
 - `last()`: returns an entry in D with the largest key.
 - `successors(k)`: returns an iterable collection of all entries in D with keys greater than or equal to k , in non-decreasing order.
 - `predecessors(k)`: returns an iterable collection of all entries in D with keys less than or equal to k , in non-increasing order.

Ordered Dictionary Implementations

- An **Ordered Dictionary** can be implemented efficiently using three main data structures:
 - A **balanced search tree**, (e.g. **red-black** or **AVL binary trees**); and a **B-tree** (or its variants) can also be used when the map is too large to reside entirely in main memory.
 - A **search table**, (e.g. **Sorted ArrayList**).
 - A **skip list**.

Comparing Ordered Dictionary Implementations

- All implementations require $O(n)$ space. Assume the following:
 - n : The number of entries in the dictionary,
 - s : The size of the collection returned by operation **findAll(k)**.
- The time requirements of all operations are shown in the following table:

Operation Using	find	findAll	remove	insert	first	last	successor	predecessor
Balanced Search Tree	$O(\log n)$	$O(\log n + s)$	$O(\log n)$		$O(\log n)$		$O(\log n)$	$O(\log n)$
Search table	$O(\log n)$	$O(\log n + s)$	$O(n)$		$O(1)$		$O(1)$	$O(1)$
Skip List	$O(\log n)$ exp.	$O(\log n + s)$ exp.	$O(\log n)$ exp.		$O(1)$		$O(\log n)$ exp.	$O(\log n)$ exp.